

Quantium VI

Rolly Mougoue

2022-04-19

Task1- Data preparation and customer analytics

Load required libraries and datasets

We'll start with setting up our working environment by installing and loading all the necessary packages (Tidyverse, Lubridate, dplyr, pad...) and setting the working directory for easy upload of our datasets.

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr   1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(data.table)
```

```
##
```

```
## Attaching package: 'data.table'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      between, first, last
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      transpose
```

```
library(ggplot2)
library(ggmosaic)
library(readr)
library(readxl)
```

```
## Setting the work directory and uploading our data sets

setwd(dir = "D:/Google Data Analytics/The Forage Projects/Quantium/Datasets")

purchase_behaviour <- read_csv("QVI_purchase_behaviour.csv")

## Rows: 72637 Columns: 3

## -- Column specification -----
## Delimiter: ","
## chr (2): LIFESTAGE, PREMIUM_CUSTOMER
## dbl (1): LYLTY_CARD_NBR
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

Transactions <- read_excel("QVI_transaction_data.xlsx")
```

Exploratory data analysis

The first step in any analysis is to first understand the data. Let's take a look at each of the datasets provided.

Examining transaction data

We can use `str()` or `glimpse()` to look at the format of each column and see a sample of the data. As we have read in the dataset as `adata.table` object, we can also run `Transactions` in the console to see a sample of the data or use `head(Transactions)` to look at the first 10 rows. Let's check if columns we would expect to be numeric are in numeric form and date columns are in date format.

```
str(Transactions)
```

```
## tibble [264,836 x 8] (S3: tbl_df/tbl/data.frame)
## $ DATE           : num [1:264836] 43390 43599 43605 43329 43330 ...
## $ STORE_NBR      : num [1:264836] 1 1 1 2 2 4 4 4 5 7 ...
## $ LYLTY_CARD_NBR: num [1:264836] 1000 1307 1343 2373 2426 ...
## $ TXN_ID         : num [1:264836] 1 348 383 974 1038 ...
## $ PROD_NBR       : num [1:264836] 5 66 61 69 108 57 16 24 42 52 ...
## $ PROD_NAME      : chr [1:264836] "Natural Chip          Compny SeaSalt175g" "CCs Nacho Cheese    175g
## $ PROD_QTY       : num [1:264836] 2 3 2 5 3 1 1 1 1 2 ...
## $ TOT_SALES      : num [1:264836] 6 6.3 2.9 15 13.8 5.1 5.7 3.6 3.9 7.2 ...
```

```
Transactions <- data.table(Transactions)
purchase_behaviour <- data.table(purchase_behaviour)
```

Examining the DATE variable

We can see that the date column is in an integer format. Let's change this to a date format.

```
#### Convert DATE column to a date format
Transactions$DATE <- as.Date(Transactions$DATE, origin = "1899-12-30")
```

Examining the PROD_NAME variable

We should check that we are looking at the right products by examining PROD_NAME.

```
#### Examine PROD_NAME
Prod_Names <- as.data.frame(table(Transactions$PROD_NAME))
```

Looks like we are definitely looking at potato chips but how can we check that these are all chips? We can do some basic text analysis by summarising the individual words in the product name.

```
#### Examine the words in PROD_NAME to see if there are any incorrect entries
productWords <- data.table(unlist(strsplit(unique(Transactions$PROD_NAME), " ")))
setnames(productWords, 'words')
```

As we are only interested in words that will tell us if the product is chips or not, let's remove all words with digits and special characters such as '&' from our set of product words. We can do this using `grepl()`.

```
#### Removing digits
productWords <- productWords[grepl("\\d", words) == FALSE, ]
```

```
#### Removing special characters
productWords <- productWords[grepl("[[:alpha:]]", words), ]
```

```
#### Let's look at the most common words by counting the number of times a word appears and sorting them
productWords <- productWords[, .N, words][order(N, decreasing = TRUE)]
```

There are salsa products in the dataset but we are only interested in the chips category, so let's remove these.

```
#### remove the salsa product
Transactions <- Transactions %>%
  dplyr::filter(!grepl("Salsa", Transactions$PROD_NAME))

#### summarizing the data
summary(Transactions)
```

##	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID
##	Min. :2018-07-01	Min. : 1.0	Min. : 1000	Min. : 1
##	1st Qu.:2018-09-30	1st Qu.: 70.0	1st Qu.: 70015	1st Qu.: 67569
##	Median :2018-12-30	Median :130.0	Median : 130367	Median : 135183
##	Mean :2018-12-30	Mean :135.1	Mean : 135531	Mean : 135131
##	3rd Qu.:2019-03-31	3rd Qu.:203.0	3rd Qu.: 203084	3rd Qu.: 202654
##	Max. :2019-06-30	Max. :272.0	Max. :2373711	Max. :2415841
##	PROD_NBR	PROD_NAME	PROD_QTY	TOT_SALES

```
## Min.      : 1.00    Length:246742    Min.      : 1.000    Min.      : 1.700
## 1st Qu.: 26.00    Class :character    1st Qu.: 2.000    1st Qu.: 5.800
## Median : 53.00    Mode  :character    Median : 2.000    Median : 7.400
## Mean    : 56.35                                Mean    : 1.908    Mean    : 7.321
## 3rd Qu.: 87.00                                3rd Qu.: 2.000    3rd Qu.: 8.800
## Max.    :114.00                                Max.    :200.000    Max.    :650.000
```

In product quantity it is visible that in PROD_QTY column the max value is of 200 i.e 200 quantity were purchased at once.. hence this can be considered as a outlier to our data.. there are no nulls as all the summary statistics have a numerical value # finding the outlier in PROD_QTY

```
#### Filter the dataset to find the outlier
filter(Transactions, PROD_QTY == "200")
```

```
##          DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
## 1: 2018-08-19      226          226000 226201        4
## 2: 2019-05-20      226          226000 226210        4
##          PROD_NAME PROD_QTY TOT_SALES
## 1: Dorito Corn Chp    Supreme 380g      200      650
## 2: Dorito Corn Chp    Supreme 380g      200      650
```

We can observe that there were two transactions done that had product quantity > 10. Both the transaction were done using the same loyalty card number and the same store. Let's see if the customer has had other transactions

```
filter(Transactions, LYLTY_CARD_NBR == 226000)
```

```
##          DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
## 1: 2018-08-19      226          226000 226201        4
## 2: 2019-05-20      226          226000 226210        4
##          PROD_NAME PROD_QTY TOT_SALES
## 1: Dorito Corn Chp    Supreme 380g      200      650
## 2: Dorito Corn Chp    Supreme 380g      200      650
```

It looks like this customer has only had the two transactions over the year and is not an ordinary retail customer. The customer might be buying chips for commercial purposes instead. We'll remove this loyalty card number from further analysis.

```
# Filter out the customer based on the loyalty card number
Transactions <- filter(Transactions, LYLTY_CARD_NBR != 226000)
#### Re-examine transaction data
summary(Transactions)
```

```
##          DATE          STORE_NBR    LYLTY_CARD_NBR    TXN_ID
## Min.      :2018-07-01    Min.      : 1.0    Min.      : 1000    Min.      : 1
## 1st Qu.:2018-09-30    1st Qu.: 70.0    1st Qu.: 70015    1st Qu.: 67569
## Median :2018-12-30    Median :130.0    Median : 130367    Median : 135182
## Mean     :2018-12-30    Mean   :135.1    Mean   : 135530    Mean   : 135130
## 3rd Qu.:2019-03-31    3rd Qu.:203.0    3rd Qu.: 203083    3rd Qu.: 202652
## Max.     :2019-06-30    Max.   :272.0    Max.   :2373711    Max.   :2415841
##          PROD_NBR    PROD_NAME    PROD_QTY    TOT_SALES
## Min.      : 1.00    Length:246740    Min.      :1.000    Min.      : 1.700
```

```
## 1st Qu.: 26.00    Class :character    1st Qu.:2.000    1st Qu.: 5.800
## Median : 53.00    Mode  :character    Median :2.000    Median : 7.400
## Mean   : 56.35                    Mean   :1.906    Mean   : 7.316
## 3rd Qu.: 87.00                    3rd Qu.:2.000    3rd Qu.: 8.800
## Max.   :114.00                    Max.    :5.000    Max.    :29.500
```

That's better. Now, let's look at the number of transaction lines over time to see if there are any obvious data issues such as missing data.

```
#### Count the number of transactions by date
```

```
Transactions %>%
  group_by(DATE) %>%
  summarise(num_trans = n())
```

```
## # A tibble: 364 x 2
##   DATE          num_trans
##   <date>          <int>
## 1 2018-07-01         663
## 2 2018-07-02         650
## 3 2018-07-03         674
## 4 2018-07-04         669
## 5 2018-07-05         660
## 6 2018-07-06         711
## 7 2018-07-07         695
## 8 2018-07-08         653
## 9 2018-07-09         692
## 10 2018-07-10        650
## # ... with 354 more rows
```

There's only 364 rows, meaning only 364 dates which indicates a missing date. Let's create a sequence of dates from 1 Jul 2018 to 30 Jun 2019 and use this to create a chart of number of transactions over time to find the missing date.

```
#### Create a sequence of dates and join this the count of transactions by date
```


```
allDates <- data.table(seq(as.Date("2018/07/01"), as.Date("2019/06/30"), by = "day"))
setnames(allDates, "DATE")
transactions_by_day <- merge(allDates, Transactions, by = "DATE", all.y = TRUE, all.x = TRUE) %>%
  group_by(DATE) %>%
  summarise(num_trans = n())
```

Setting plot themes to format graphs

```
theme_set(theme_bw())
theme_update(plot.title = element_text(hjust = 0.5))
```


Plot transactions over time

```
ggplot(transactions_by_day, aes(x = DATE, y = num_trans)) +
  geom_line(col = "Blue") +
  labs(x = "Date", y = "Number of transactions", title = "Transactions over time") +
  scale_x_date(breaks = "1 month") + ylim(600, 900) +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```



We can see that there is an increase in purchases in December and a break in late December. Let's zoom in on this.

```
#### Filter to December and look at individual days
ggplot(transactions_by_day[transactions_by_day$DATE > "2018-12-01" & transactions_by_day$DATE < "2018-12-31",]) +
  geom_line(aes(x = DATE, y = num_trans)) +
  labs(x = "Day", y = "Number of transactions", title = "December Transactions over time") +
  scale_x_date(breaks = "1 day") + ylim(600, 900) +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```



We can see that the increase in sales occurs in the lead-up to Christmas and that there are zero sales on Christmas day itself. This is due to shops being closed on Christmas day.

Examining the Chips Pack Sizes and Brand

Now that we are satisfied that the data no longer has outliers, we can move on to creating other features such as brand of chips or pack size from `PROD_NAME`. We will start with pack size.

```
#### We can work this out by taking the digits that are in PROD_NAME
Transactions[, PACK_SIZE := parse_number(PROD_NAME)]
```

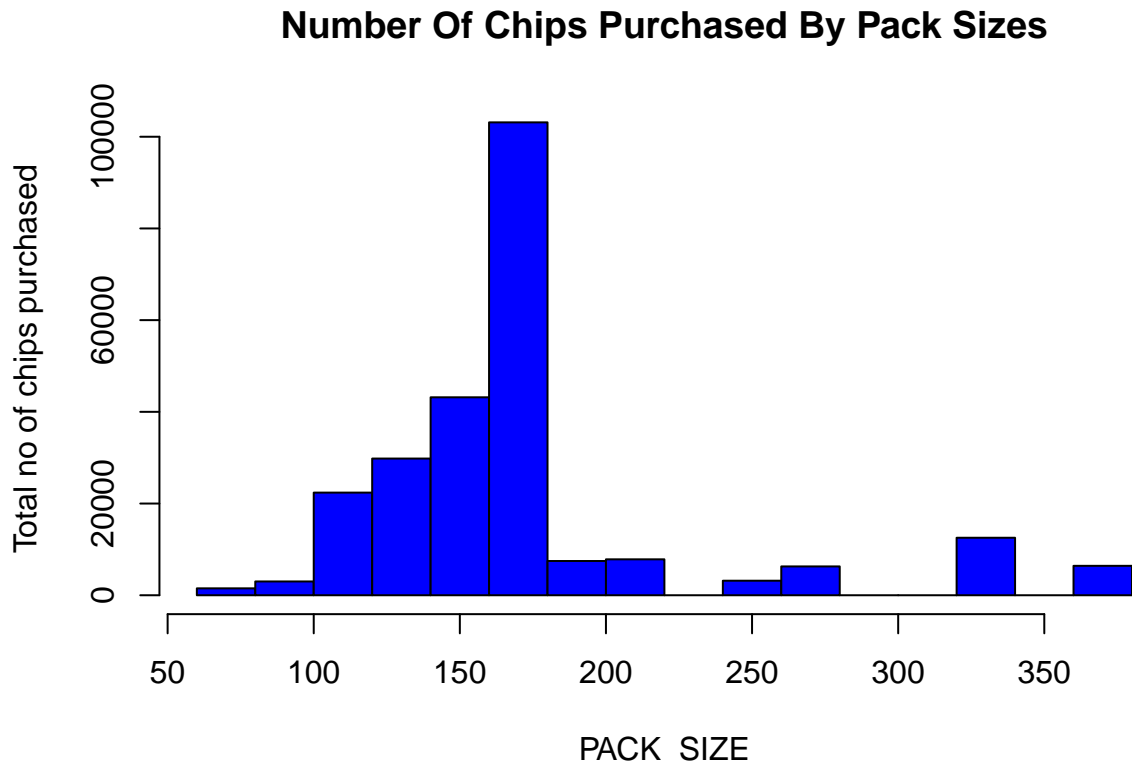
```
## Warning in `[.data.table'(Transactions, , `:=`(PACK_SIZE,
## parse_number(PROD_NAME)))': Invalid .internal.selfref detected and fixed by
## taking a (shallow) copy of the data.table so that := can add this new column
## by reference. At an earlier point, this data.table has been copied by R (or
## was created manually using structure() or similar). Avoid names<- and attr<-
## which in R currently (and oddly) may copy the whole data.table. Use set* syntax
## instead to avoid copying: ?set, ?setnames and ?setattr. If this message doesn't
## help, please report your use case to the data.table issue tracker so the root
## cause can be fixed or this message improved.
```

```
#### Let's check if the pack sizes look sensible
Transactions[, .N, PACK_SIZE][order(PACK_SIZE)]
```

##	PACK_SIZE	N
## 1:	70	1507
## 2:	90	3008
## 3:	110	22387
## 4:	125	1454
## 5:	134	25102
## 6:	135	3257
## 7:	150	40203
## 8:	160	2970
## 9:	165	15297
## 10:	170	19983
## 11:	175	66390
## 12:	180	1468
## 13:	190	2995
## 14:	200	4473
## 15:	210	6272
## 16:	220	1564
## 17:	250	3169
## 18:	270	6285
## 19:	330	12540
## 20:	380	6416

The largest size is 380g and the smallest size is 70g - seems sensible! Now Let's plot a histogram of PACK_SIZE since we know that it is a categorical variable and not a continuous variable even though it is numeric.

```
#### plotting histogram of the packsize
options(scipen=999) # turn off scientific notations like 1e+05
hist(Transactions[, PACK_SIZE], col = "blue",border = "black" , xlab = "PACK SIZE", ylab = "Total no o
```

the plot looks reasonable with no outliers and from the plot it can be seen that the the packs of size 170-180 was purchased the most. Now to create brands, we can use the first word in PROD_NAME to work out the brand name...

```
#### we'll use the first word of the PROD_NAME to create our data of brand
Transactions[, Brands := toupper(substr(PROD_NAME, 1, regexr(pattern = ' ', PROD_NAME) - 1))]
```

```
#### Checking brands
Transactions[, .N, by = Brands][order(-N)]
```

```
##      Brands      N
##  1:  KETTLE 41288
##  2:  SMITHS 27390
##  3: PRINGLES 25102
##  4:  DORITOS 22041
##  5:   THINS 14075
##  6:    RRD 11894
##  7: INFUZIONI 11057
##  8:    WW 10320
##  9:   COBS  9693
## 10: TOSTITOS  9471
## 11: TWISTIES  9454
## 12: TYRRELLS  6442
## 13:   GRAIN  6272
## 14:  NATURAL  6050
```

```
## 15: CHEEZELS 4603
## 16: CCS 4551
## 17: RED 4427
## 18: DORITO 3183
## 19: INFZNS 3144
## 20: SMITH 2963
## 21: CHEETOS 2927
## 22: SNBTS 1576
## 23: BURGER 1564
## 24: WOOLWORTHS 1516
## 25: GRNWVES 1468
## 26: SUNBITES 1432
## 27: NCC 1419
## 28: FRENCH 1418
## Brands N
```

Some of the brand names look like they are of the same brands - such as RED and RRD, which are both Red Rock Deli chips. Let's combine these rows together.

```
Transactions[Brands == "RED", Brands := "RRD"]
Transactions[Brands == "SNBTS", Brands := "SUNBITES"]
Transactions[Brands == "INFZNS", Brands := "INFUZIONI"]
Transactions[Brands == "WW", Brands := "WOOLWORTHS"]
Transactions[Brands == "SMITH", Brands := "SMITHS"]
Transactions[Brands == "NCC", Brands := "NATURAL"]
Transactions[Brands == "DORITO", Brands := "DORITOS"]
Transactions[Brands == "GRAIN", Brands := "GRNWVES"]
```

Checking again brand names

```
Transactions[, .N, by = Brands][order(Brands)]
```

```
## Brands N
## 1: BURGER 1564
## 2: CCS 4551
## 3: CHEETOS 2927
## 4: CHEEZELS 4603
## 5: COBS 9693
## 6: DORITOS 25224
## 7: FRENCH 1418
## 8: GRNWVES 7740
## 9: INFUZIONI 14201
## 10: KETTLE 41288
## 11: NATURAL 7469
## 12: PRINGLES 25102
## 13: RRD 16321
## 14: SMITHS 30353
## 15: SUNBITES 3008
## 16: THINS 14075
## 17: TOSTITOS 9471
## 18: TWISTIES 9454
## 19: TYRRELLS 6442
## 20: WOOLWORTHS 11836
```

Now 8 of our rows that had similar brand has been merged and now we can finally stop our data exploration and continue further..

Checking the customer data

Examining customer data

Now that we are happy with the transaction dataset, let's have a look at the customer dataset.

```
str(purchase_behaviour)
```

```
## Classes 'data.table' and 'data.frame': 72637 obs. of 3 variables:
## $ LYLTY_CARD_NBR : num 1000 1002 1003 1004 1005 ...
## $ LIFESTAGE : chr "YOUNG SINGLES/COUPLES" "YOUNG SINGLES/COUPLES" "YOUNG FAMILIES" "OLDER SI
## $ PREMIUM_CUSTOMER: chr "Premium" "Mainstream" "Budget" "Mainstream" ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
summary(purchase_behaviour)
```

```
## LYLTY_CARD_NBR LIFESTAGE PREMIUM_CUSTOMER
## Min. : 1000 Length:72637 Length:72637
## 1st Qu.: 66202 Class :character Class :character
## Median : 134040 Mode :character Mode :character
## Mean : 136186
## 3rd Qu.: 203375
## Max. : 2373711
```

we can see that the data is mostly descriptive. It gives the description of the customer who purchased the chip, the loyalty card number is a numeric vector while lifestage and premium_customer are character vectors

Let's have a closer look at the LIFESTAGE and PREMIUM_CUSTOMER columns.

```
##### Examining the values of lifestage and premium_customer
purchase_behaviour[, .N, by = LIFESTAGE][order(-N)]
```

```
##          LIFESTAGE      N
## 1:      RETIREES 14805
## 2:  OLDER SINGLES/COUPLES 14609
## 3:  YOUNG SINGLES/COUPLES 14441
## 4:      OLDER FAMILIES 9780
## 5:      YOUNG FAMILIES 9178
## 6:  MIDAGE SINGLES/COUPLES 7275
## 7:      NEW FAMILIES 2549
```

```
purchase_behaviour[, .N, by = PREMIUM_CUSTOMER][order(-N)]
```

```
##    PREMIUM_CUSTOMER      N
## 1:      Mainstream 29245
## 2:          Budget 24470
## 3:          Premium 18922
```

Merge transaction data to customer data

```
data <- merge(Transactions, purchase_behaviour , all.x = TRUE)
```

As the number of rows in `data` is the same as that of `Transactions`, we can be sure that no duplicates were created. This is because we created `data` by setting `all.x = TRUE` (in other words, a left join) which means take all the rows in `Transactions` and find rows with matching values in shared columns and then joining the details in these rows to the `x` or the first mentioned table.

Let's also check if some customers were not matched on by checking for nulls.

```
colSums(is.na(data))
```

```
##  LYLTY_CARD_NBR      DATE      STORE_NBR      TXN_ID
##           0           0           0           0
##      PROD_NBR      PROD_NAME      PROD_QTY      TOT_SALES
##           0           0           0           0
##      PACK_SIZE      Brands      LIFESTAGE      PREMIUM_CUSTOMER
##           0           0           0           0
```

Great, there are no nulls! So all our customers in the transaction data has been accounted for in the customer dataset. Hence we can proceed further since all the data was matched properly

```
#### saving the cleaned file for further analysis
write.csv(data, "QVI_Cleaned_data.csv")
```

We've now completed data exploration from different datasets.

Data analysis on customer segments

Now that the data is ready for analysis, we can define some metrics of interest to the client such as: Who spends the most on chips (total sales), describing customers by lifestage and how premium their general purchasing behaviour is - How many customers are in each segment - How many chips are bought per customer by segment - What's the average chip price by customer segment

Let's start with calculating total sales by `LIFESTAGE` and `PREMIUM_CUSTOMER` and plotting the split by these segments to describe which customer segment contribute most to chip sales.

```
# Total sales by LIFESTAGE and PREMIUM_CUSTOMER
Total_Sales <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(Sales = sum(TOT_SALES, na.rm = TRUE)) %>%
  na.omit()
```

```
## 'summarise()' has grouped output by 'LIFESTAGE'. You can override using the
## '.groups' argument.
```

```
# create plot
plot <- ggplot(data = Total_Sales) +
  geom_mosaic(aes(weight = Sales, x = product(PREMIUM_CUSTOMER, LIFESTAGE) , fill = PREMIUM_CUSTOMER)) +
  labs(x = "Lifestage", y = "Premium customer", title = "Proportion of sales By LifeStage/Customer") +
```

```
# Plot and label with proportion of sales
```

```
plot +
```


```
  geom_text(data = ggplot_build(plot)$data[[1]], aes(x = (xmin + xmax)/2 , y = (ymin + ymax)/2, label =
```

```
## Warning: 'unite_()' was deprecated in tidyr 1.2.0.
```

```
## Please use 'unite()' instead.
```

```
## This warning is displayed once every 8 hours.
```

```
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was generated.
```



We can observe from the plot that mainstream- young/single couples and mainstream retirees contribute the most to there being higher sales of chips which isn't the case for budget- older families segment.

```
## Number of customers by LIFESTAGE and PREMIUM_CUSTOMER
```

```
Customers <- data[, .(customers = uniqueN(LYLT_CARD_NBR)), .(LIFESTAGE, PREMIUM_CUSTOMER)] [order(-custo
```

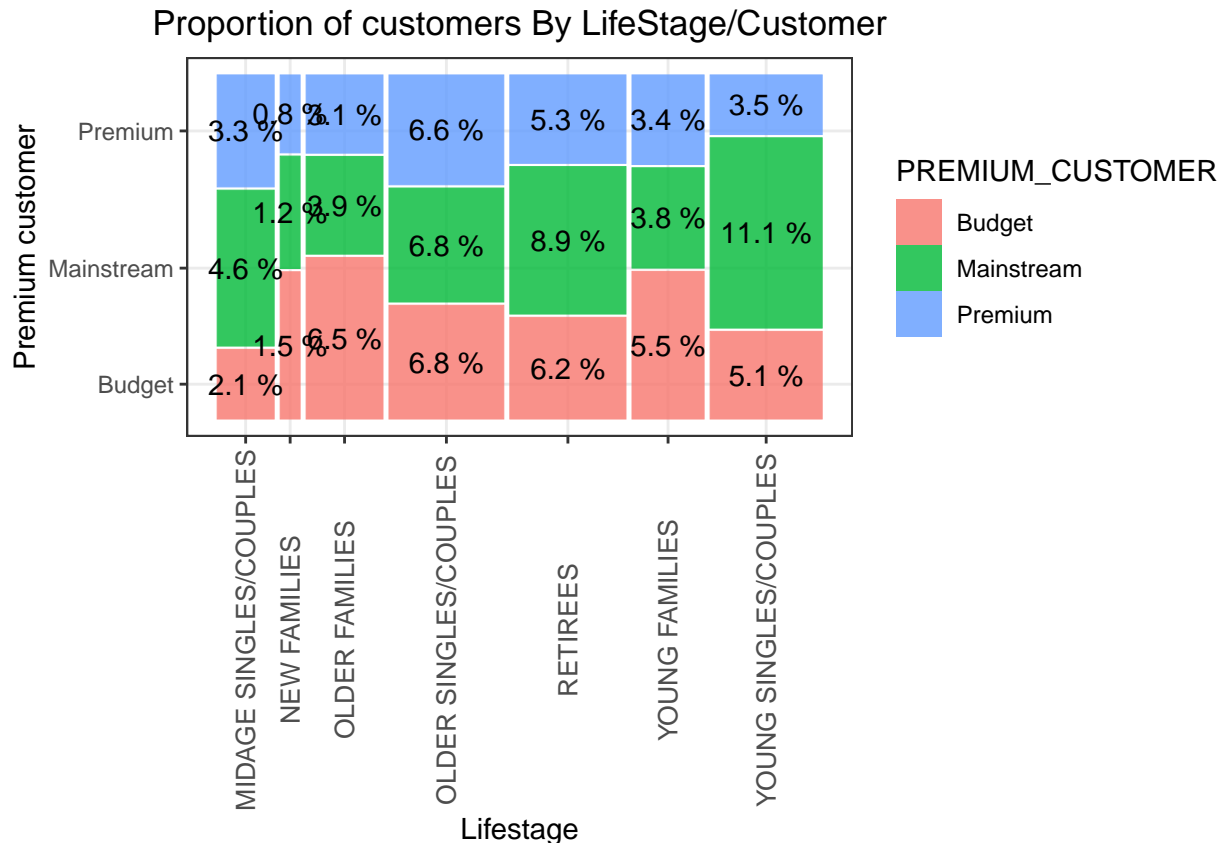
```
# Create plot
```

```
Plot <- ggplot(data = Customers) +
```

```
  geom_mosaic(aes(weight = customers, x = product(PREMIUM_CUSTOMER, LIFESTAGE), fill = PREMIUM_CUSTOMER,
```

```
  labs(x = "Lifestage", y = "Premium customer", title = "Proportion of customers By LifeStage/Customer").
```

```
Plot + geom_text(data = ggplot_build(Plot)$data[[1]], aes(x = (xmin + xmax)/2 , y = (ymin + ymax)/2, lab
```



There are more Mainstream - young singles/couples and Mainstream - retirees who buy chips. This contributes to there being more sales to these customer segments but this is not a major driver for the Budget-Older families segment.

Higher sales may also be driven by number of chips bought by each customer.. hence we'll try to plot average number of chips(average number of PROD_QTY) by lifestage and premium_customer. Let's have a look at this next

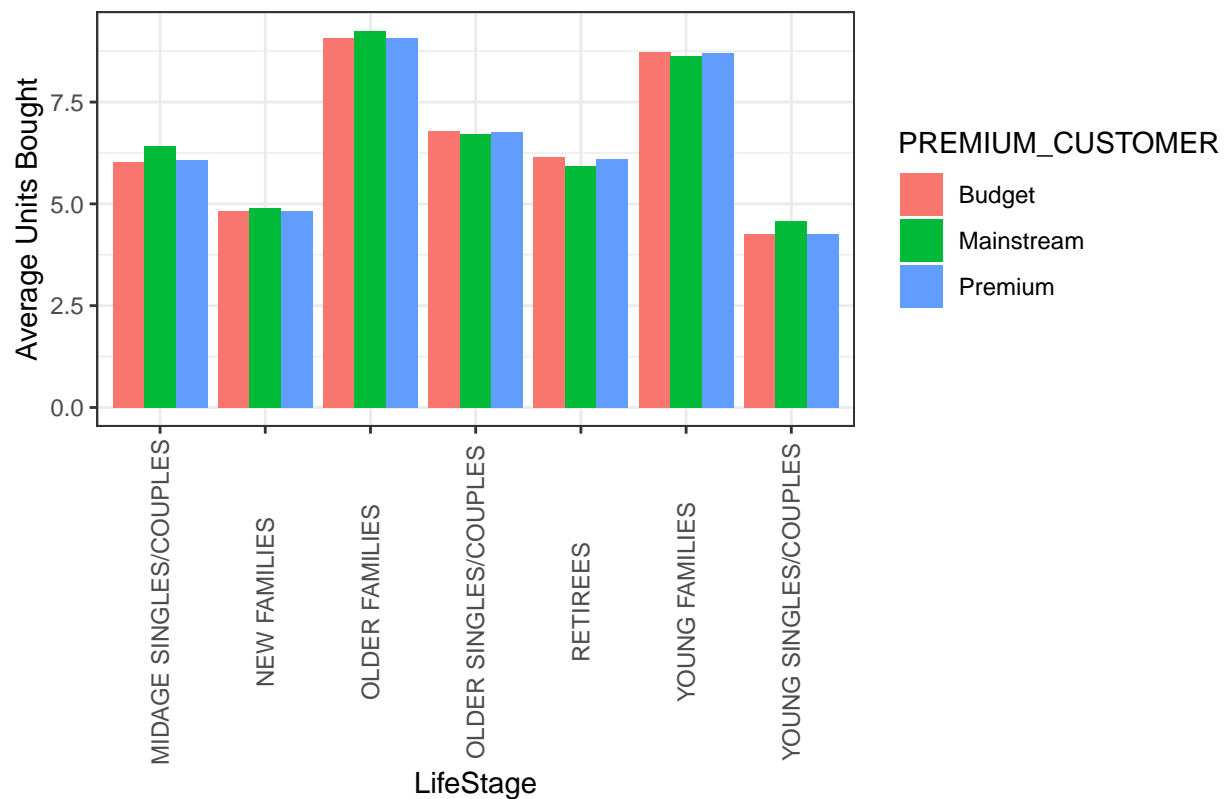
```
#### Average number of units per customer by LIFESTAGE and PREMIUM_CUSTOMER
```

```
Average_Units <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(Quantity = sum(PROD_QTY)/uniqueN(LYLT_CARD_NBR)) %>%
  na.omit()
```

```
## 'summarise()' has grouped output by 'LIFESTAGE'. You can override using the
## '.groups' argument.
```

```
ggplot(data = Average_Units, aes(weight = Quantity , x = LIFESTAGE, fill = PREMIUM_CUSTOMER)) +
  geom_bar(position = position_dodge()) +
  labs(title = " Average units per Customers by LIFESTAGE by Categorie", x = "LifeStage", y = "Average")
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```

Average units per Customers by LIFESTAGE by Categorie



Older families and young families in general buy more chips per customer Let's also investigate the average price per unit chips bought for each customer segment as this is also a driver of total sales.


Average price per unit by LIFESTAGE and PREMIUM_CUSTOMER

```
avg_price <- data[, .(AVG = sum(TOT_SALES)/sum(PROD_QTY)), .(LIFESTAGE, PREMIUM_CUSTOMER)][order(-AVG)]
avg_price <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(AVG = sum(TOT_SALES)/sum(PROD_QTY)) %>%
  na.omit()
```

'summarise()' has grouped output by 'LIFESTAGE'. You can override using the
'.groups' argument.

Create plot

```
ggplot(data = avg_price, aes(weight = AVG, x = LIFESTAGE, fill = PREMIUM_CUSTOMER)) + geom_bar(position
```



Mainstream midage and young singles and couples are more willing to pay more per packet of chips compared to their budget and premium counterparts. This may be due to premium shoppers being more likely to buy healthy snacks and when they buy chips, this is mainly for entertainment purposes rather than their own consumption. This is also supported by there being fewer premium midage and young singles and couples buying chips compared to their mainstream counterparts.

As the difference between the average price per unit is not same we can check if this difference is statically significant or not by performing independent t-test between mainstream vs premium and budget midage and young young single couples.

```
#### young singles and couples
pricePerUnit <- data[, price := TOT_SALES/PROD_QTY]
t.test(data[LIFESTAGE %in% c("YOUNG SINGLES/COUPLES", "MIDAGE SINGLES/COUPLES") & PREMIUM_CUSTOMER == "Mainstream", data[LIFESTAGE %in% c("YOUNG SINGLES/COUPLES", "MIDAGE SINGLES/COUPLES") & PREMIUM_CUSTOMER != "Mainstream"], alternative = "greater")

##
## Welch Two Sample t-test
##
## data: data[LIFESTAGE %in% c("YOUNG SINGLES/COUPLES", "MIDAGE SINGLES/COUPLES") & PREMIUM_CUSTOMER == "Mainstream", data[LIFESTAGE %in% c("YOUNG SINGLES/COUPLES", "MIDAGE SINGLES/COUPLES") & PREMIUM_CUSTOMER != "Mainstream"], alternative = "greater"]
## t = 37.624, df = 54791, p-value < 0.00000000000000022
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## 0.3187234 Inf
## sample estimates:
## mean of x mean of y
## 4.039786 3.706491
```


The t-test results in a p-value of 2.2e-16. Meaning the unit price for mainstream, young and mid-age singles and couples ARE significantly higher than that of budget or premium, young and midage singles and couples.

Deep dive into specific customer segments for insights We have found quite a few interesting insights that we can dive deeper into. We might want to target customer segments that contribute the most to sales to retain them or further increase sales. Let's look at Mainstream - young singles/couples. For instance, let's find out if they tend to buy a particular brand of chips.

```
# Deep dive into Mainstream, young singles/couples
Customer_Type1 <- data[LIFESTAGE == "YOUNG SINGLES/COUPLES" & PREMIUM_CUSTOMER == "Mainstream",]
Customer_Type2<- data[!(LIFESTAGE == "YOUNG SINGLES/COUPLES" & PREMIUM_CUSTOMER == "Mainstream"),]

## Brand affinity compared to the rest of the population
quantity_type1 <- Customer_Type1[, sum(PROD_QTY)]

quantity_type2 <- Customer_Type2[, sum(PROD_QTY)]

quantity_type1_by_brand <- Customer_Type1[, .(targetSegment = sum(PROD_QTY)/quantity_type1), by = Brands]
quantity_type2_by_brand <- Customer_Type2[, .(other = sum(PROD_QTY)/quantity_type2), by = Brands]

brand_proportions <- merge(quantity_type1_by_brand, quantity_type2_by_brand)[, affinityToBrand := targetSegment/other]

brand_proportions[order(-affinityToBrand)]
```

##	Brands	targetSegment	other	affinityToBrand
## 1:	TYRRELLS	0.031552795	0.025692464	1.2280953
## 2:	TWISTIES	0.046183575	0.037876520	1.2193194
## 3:	DORITOS	0.122760524	0.101074684	1.2145526
## 4:	KETTLE	0.197984817	0.165553442	1.1958967
## 5:	TOSTITOS	0.045410628	0.037977861	1.1957131
## 6:	PRINGLES	0.119420290	0.100634769	1.1866703
## 7:	COBS	0.044637681	0.039048861	1.1431238
## 8:	INFUZIONS	0.064679089	0.057064679	1.1334347
## 9:	THINS	0.060372671	0.056986370	1.0594230
## 10:	GRNWVES	0.032712215	0.031187957	1.0488733
## 11:	CHEEZELS	0.017971014	0.018646902	0.9637534
## 12:	SMITHS	0.096369910	0.124583692	0.7735355
## 13:	FRENCH	0.003947550	0.005758060	0.6855694
## 14:	CHEETOS	0.008033126	0.012066591	0.6657329
## 15:	RRD	0.043809524	0.067493678	0.6490908
## 16:	NATURAL	0.019599724	0.030853989	0.6352412
## 17:	CCS	0.011180124	0.018895650	0.5916771
## 18:	SUNBITES	0.006349206	0.012580210	0.5046980
## 19:	WOOLWORTHS	0.024099379	0.049427188	0.4875733
## 20:	BURGER	0.002926156	0.006596434	0.4435967

We can see that : - Mainstream young singles/couples are 23% more likely to purchase Tyrrells chips compared to the rest of the population - Mainstream young singles/couples are 56% less likely to purchase Burger Rings compared to the rest of the population

Let's also find out if our target segment tends to buy larger packs of chips.

```
# Preferred pack size compared to the rest of the population
quantity_type1_by_pack <- Customer_Type1[, .(targetSegment = sum(PROD_QTY)/quantity_type1), by = PACK_SIZE]
quantity_type2_by_pack <- Customer_Type2[, .(other = sum(PROD_QTY)/quantity_type2), by = PACK_SIZE]

pack_proportions <- merge(quantity_type1_by_pack, quantity_type2_by_pack)[, affinityToPack := targetSegment]

pack_proportions[order(-affinityToPack)]
```

##	PACK_SIZE	targetSegment	other	affinityToPack
## 1:	270	0.031828847	0.025095929	1.2682873
## 2:	380	0.032160110	0.025584213	1.2570295
## 3:	330	0.061283644	0.050161917	1.2217166
## 4:	134	0.119420290	0.100634769	1.1866703
## 5:	110	0.106280193	0.089791190	1.1836372
## 6:	210	0.029123533	0.025121265	1.1593180
## 7:	135	0.014768806	0.013075403	1.1295106
## 8:	250	0.014354727	0.012780590	1.1231662
## 9:	170	0.080772947	0.080985964	0.9973697
## 10:	150	0.157598344	0.163420656	0.9643722
## 11:	175	0.254989648	0.270006956	0.9443818
## 12:	165	0.055652174	0.062267662	0.8937572
## 13:	190	0.007481021	0.012442016	0.6012708
## 14:	180	0.003588682	0.006066692	0.5915385
## 15:	160	0.006404417	0.012372920	0.5176157
## 16:	90	0.006349206	0.012580210	0.5046980
## 17:	125	0.003008972	0.006036750	0.4984423
## 18:	200	0.008971705	0.018656115	0.4808989
## 19:	70	0.003036577	0.006322350	0.4802924
## 20:	220	0.002926156	0.006596434	0.4435967

It looks like Mainstream young singles/couples are 27% more likely to purchase a 270g pack of chips compared to the rest of the population but let's dive into what brands sell this pack size.

```
data[PACK_SIZE == 270, unique(PROD_NAME)]
```

```
## [1] "Twisties Cheese      270g" "Twisties Chicken270g"
```

The only brand offering 270g packs are Twisties and this may instead be reflecting a higher likelihood of purchasing them.

Conclusion // Insights

Let's recap what we've found during this first analysis!

Sales have mainly been due to Budget - older families, Mainstream - young singles/couples, and Mainstream retirees shoppers. We noticed a correlation between high spend in chips for mainstream young singles/couples and retirees and their number in the sample as compared to other buyers. Mainstream, midage and young singles and couples are also more likely to pay more per packet of chips. This is indicative of impulse buying behaviour. ****We've also found that Mainstream young singles and couples are 23% more likely to purchase Tyrrells chips compared to the rest of the population**